

Algorithmic Bias in Educational Technology and Student Surveillance

The Issue

AI-powered learning platforms, remote proctoring software, and predictive analytics systems are now deeply embedded in educational institutions, spanning from K-12 to higher education. While they are often marketed as objective, efficient, and innovative solutions to longstanding educational challenges, these technologies can unintentionally magnify existing inequalities, perpetuate systemic discrimination, and infringe on student privacy in novel and troubling ways. The so-called "black box" nature of many algorithms means that bias is not only present but also invisible and unaccountable. Students, parents, and even educators often have little understanding of—or control over—how vast amounts of personal and academic data are collected, processed, and used to make critical decisions. These decisions, which may include admissions, academic placements, and disciplinary actions, can profoundly shape students' educational experiences and long-term opportunities. Moreover, the increasing reliance on algorithmic systems risks shifting power away from human judgment and eroding trust between students and educational institutions.

Why It Matters

Algorithmic Bias: The Illusion of Objectivity

- Algorithms used for college admissions, course recommendations, or "at-risk" student identification can perpetuate racial and socioeconomic bias. While presented as neutral and data-driven, these systems are trained on historical data that reflects existing inequalities—meaning they learn to replicate discriminatory patterns.
- Admissions algorithms may disadvantage applicants from underrepresented backgrounds. When AI systems evaluate applications, they may weigh factors like high school prestige, zip code, extracurricular activities, or writing style in ways that systematically favor wealthy, white applicants. Students from under-resourced schools or neighborhoods get algorithmically sorted out before human eyes ever see their applications.
- "At-risk" student identification systems can create self-fulfilling prophecies. Predictive analytics flag students as likely to fail or drop out based on factors like demographics, prior test scores, attendance, or engagement metrics. Once flagged, students may receive interventions that are stigmatizing, lower expectations from teachers, or reduced access to advanced courses—reinforcing the very outcomes the algorithm predicted.
- Course recommendation algorithms track students into paths that limit opportunity. AI systems may steer students toward remedial courses, vocational tracks, or "easier" options based on biased assessments of their potential, perpetuating educational stratification along race and class lines.

- Grading and assessment algorithms may penalize diverse communication styles. Automated essay scoring systems often favor formulaic, standardized academic English, potentially marking down students who use African American Vernacular English (AAVE), code-switching, creative expression, or non-Western rhetorical structures—even when ideas are strong and clear.
- AI tutoring and adaptive learning platforms may provide lower-quality instruction to certain students. If algorithms determine students are "struggling," they may be routed to simpler content, repetitive drill-and-practice exercises, or slower pacing—denying them access to challenging, enriching material that could help them grow.
- The bias is often invisible and unaccountable. Proprietary algorithms are treated as trade secrets, so schools, teachers, students, and families cannot examine how decisions are made, challenge biased outcomes, or demand transparency.

Remote Proctoring: Surveillance Disguised as Academic Integrity

- Remote proctoring software disproportionately flags students of color and monitors students in invasive ways (scanning rooms, tracking eye movements). These systems use facial recognition, eye-tracking, audio monitoring, and behavioral analysis to detect "cheating," but the technology is rife with problems.
- Facial recognition technology fails more often for people with darker skin tones, a well-documented problem across AI systems. Black, Latino, Asian, and Indigenous students experience higher rates of false flags, technical failures, or inability to log in—causing exam anxiety, delays, accusations, or outright exam failure through no fault of their own. (Chhua et al., 2024)
- Students must show their faces on camera for extended periods, which may be uncomfortable or impossible for students who wear religious head coverings, have facial differences or disabilities, or experience gender dysphoria.
- Eye-tracking algorithms flag "suspicious" behavior like looking away from the screen, but students with ADHD, autism, anxiety, or visual processing differences may naturally move their eyes differently. Neurotypical assumptions are coded into what counts as "normal" test-taking behavior.
- Room scans invade students' private spaces. Proctoring software often requires students to pan their webcam around the room to prove no one else is present and no unauthorized materials are visible. This forces students to reveal their living conditions—small apartments, shared bedrooms, family members in the background, lack of furniture, visible poverty—to a remote proctor and potentially recorded permanently.
- Background noise and interruptions are flagged as suspicious. Students living in multigenerational households, with roommates, in noisy neighborhoods, or with young siblings may have their exams interrupted or flagged for sounds beyond their control.
- Low-income students may lack private, quiet spaces for monitored exams, further disadvantaging them. A student taking an exam from a car, a library parking lot, or a crowded kitchen is already stressed and may be penalized by proctoring algorithms that expect ideal testing environments.

- Invasive monitoring includes keystroke analysis, browser lockdowns, and audio recording. Some systems prevent students from accessing any other applications, read clipboard data, or continuously record audio "just in case." Students have been flagged for using calculators on their phone (when permitted), talking to themselves (a common study technique), or standing up to stretch.
- Students with disabilities face particular barriers. Those who need screen readers, alternative input devices, movement breaks, or testing accommodations may find proctoring software incompatible with their needs or may be flagged for "unusual" behavior that is actually their accommodation.
- The psychological toll is significant. Being watched continuously creates anxiety that impairs performance. Students report feeling dehumanized, violated, and distrusted. Some experience panic attacks during proctored exams. (ACLU RESEARCH REPORT, n.d.)
- False accusations of cheating have serious consequences: failing the exam, failing the course, academic integrity violations on permanent records, suspension, or expulsion. Appealing these decisions is often difficult, especially when the "evidence" is algorithmic scores rather than clear proof.

Data Collection and Privacy Violations

- Data collected on students—from browsing history to biometric data—raises privacy concerns and can be sold to third parties. Educational technology companies gather massive amounts of data on students, often with minimal oversight or informed consent.
- The scope of data collection is staggering: names, birthdates, addresses, biometric identifiers (face scans, voice prints, keystroke patterns), academic records, behavioral data (time on task, clicks, mistakes), social-emotional data (mood check-ins, survey responses), health information, disciplinary records, communications, location data, and more. (Nhando, 2026)
- Students may not understand how their data is used or have control over it. Privacy policies are long, legalistic, and rarely read. Students (and parents) may not realize they're consenting to extensive surveillance and data mining when they click "agree" to use a learning platform.
- Third-party vendors have inconsistent privacy protections. Educational data may be stored on servers with weak security, shared with advertisers, sold to data brokers, or used to build behavioral profiles that follow students beyond school.
- FERPA (Family Educational Rights and Privacy Act) protections have loopholes. While FERPA regulates educational records, it has broad "school official" exceptions that allow third-party companies to access student data if they provide services to schools. Once data leaves the school's control, protections weaken. (FERPA exceptions, n.d.)
- Student data can be breached or hacked. Educational technology companies are frequent targets of cyberattacks. Breaches can expose sensitive information about minors, including their locations, routines, and vulnerabilities. (FTC Takes Action Against Education Technology Provider for Failing to Secure Students' Personal Data, 2025)

- Data persistence is a concern. Information collected in elementary school may follow students through college and into employment. Behavioral flags, "risk scores," or disciplinary data could haunt students for years, affecting opportunities in ways they never know about.
- Predictive profiles can limit futures. Data analytics companies build psychological and behavioral profiles of students that may be used by colleges, employers, or others to make decisions—turning childhood mistakes or struggles into permanent digital scarlet letters.

Surveillance Culture and Its Impact on Learning

- Constant monitoring changes student behavior in harmful ways. When students know they're being watched and assessed at all times, they may become risk-averse, less creative, and less willing to experiment or make mistakes—all essential to learning.
- Trust between students and institutions erodes. Surveillance communicates that students are presumed cheaters, threats, or problems to be managed rather than learners to be supported. This undermines the educational relationship.
- Surveillance disproportionately impacts already-marginalized students. Black and Brown students, students with disabilities, LGBTQ+ students, and low-income students are most likely to be surveilled, flagged, and punished. This reinforces school-to-prison pipeline dynamics and systemic oppression.
- Students internalize surveillance. Growing up under constant digital monitoring normalizes invasive tracking, teaching students that privacy is not a right and that they should accept surveillance as a condition of participation in society.
- Critical thinking and dissent are chilled. Students may self-censor in online discussions, avoid controversial topics, or hesitate to challenge ideas if they know their participation is tracked, analyzed, and potentially used against them.

The School-to-Prison Pipeline Goes Digital

- Behavioral monitoring algorithms can flag students for "concerning" online activity, such as researching mental health, LGBTQ+ identity, protest movements, or even assigned coursework on topics like terrorism or revolution. These flags may be reported to school administrators, parents, or law enforcement. (Schools use AI to monitor kids, hoping to prevent violence. Our investigation found security risks, 2025)
- Students of color are disproportionately flagged for the same behaviors that white students engage in without consequence, mirroring disparities in school discipline more broadly. (Brislin et al., 2024, pp. 1225-1238)
- Law enforcement partnerships with ed-tech companies create pathways from classroom to criminal justice system. Some schools use software that directly notifies police of flagged behavior, turning educational surveillance into criminal surveillance.
- Students have been visited by police, interrogated, or even arrested based on algorithmic flags—sometimes for creative writing assignments, song lyrics, or venting frustration online.

Economic and Educational Inequities

- Wealthy schools and students benefit most from educational technology while bearing fewer burdens. They access premium platforms, have IT support, and can opt out of invasive tools by switching to private schools or hiring tutors.
- Low-income schools are more likely to rely on free or low-cost ed-tech that monetizes student data as part of its business model. Students in under-resourced schools thus face both more surveillance and worse technology. (Prothero, 2024)
- The digital divide intersects with surveillance. Students with limited internet or devices may struggle to comply with proctoring requirements, while also facing algorithmic penalties for inconsistent engagement with online platforms.
- "Personalized learning" can become "personalized inequality." Adaptive algorithms that sort students into differentiated learning paths may entrench rather than address achievement gaps, giving some students rigorous challenges while others get watered-down content.

Lack of Regulation and Accountability

- Educational technology is largely unregulated, with regulatory frameworks lagging far behind the rapid pace of technological change. Unlike medical devices or financial services, ed-tech faces minimal oversight regarding efficacy, safety, bias, or privacy. Companies can deploy experimental, unproven algorithms on students without demonstrating that these systems are effective, fair, or safe. There are few requirements for independent validation or public reporting of outcomes.
- Schools often lack the expertise and resources to critically evaluate ed-tech claims. Administrators, teachers, and school boards may not have the technical knowledge or time to assess the presence of algorithmic bias, data security vulnerabilities, or privacy implications. This lack of capacity leaves schools reliant on vendor marketing materials and sales pitches, which may overstate benefits and minimize risks.
- Procurement processes frequently prioritize cost, convenience, and speed over student welfare. Cash-strapped districts may adopt ed-tech tools simply because they are cheap or free, sometimes as a condition for access to basic digital infrastructure. These decisions are often made without thorough vetting of privacy policies, ongoing maintenance needs, or consideration of long-term ethical implications. Vendors may bundle products or offer incentives that make it difficult for schools to opt out of problematic systems.
- Students and families have little recourse when algorithms make mistakes, reinforce bias, or violate privacy. Challenging these decisions is often difficult or impossible. Companies may hide behind proprietary claims, preventing external audits or appeals. Schools themselves may lack clear mechanisms for students and families to contest algorithmic decisions, seek remedies, or have data deleted.
- Audit and transparency requirements remain minimal or nonexistent. Few jurisdictions mandate algorithmic impact assessments, independent bias audits, or public disclosure of how ed-tech systems make decisions and what data they use. (Algorithmic Accountability Act of 2025, n.d.)

As a result, potentially harmful systems can operate unchecked, and meaningful accountability is rare.

Resistance and Harm to Marginalized Communities

- Indigenous students' data sovereignty is violated. Tribal nations have specific concerns about data collection on Indigenous children, which may be used without consent or cultural understanding, and which violates principles of Indigenous data sovereignty.
- Undocumented students and mixed-status families face unique risks. Data collection that includes addresses, family information, or immigration status could be shared with immigration enforcement, deliberately or through breaches.
- Students with disabilities face both exclusion and over-surveillance. Inaccessible platforms exclude them from learning, while special education data is often shared extensively among staff and systems, violating privacy.
- LGBTQ+ students risk being outed. Algorithms that analyze writing, social patterns, or search history could identify students exploring their identity before they're ready to disclose, potentially putting them at risk in unsupportive homes or communities.

The Pedagogy of Surveillance vs. The Pedagogy of Trust

- Surveillance-based education assumes students are problems to manage rather than humans to nurture. It prioritizes compliance, control, and risk mitigation over curiosity, growth, and agency.
- Trust-based education values relationships, autonomy, and intrinsic motivation. It recognizes that learning requires safety to take risks, make mistakes, and explore ideas without fear of permanent consequences.
- The presence of surveillance changes what education means. When every click, glance, and word is monitored and analyzed, education becomes performance for algorithms rather than authentic intellectual and personal development.

What Students and Families Deserve

- **Transparency**
 - Clear, accessible information about what data is collected on students.
 - Explanation of how data is used, who has access, and how long it is retained.
 - Regular reporting to students and families about data practices.
- **Consent**
 - Meaningful opt-in (not just opt-out) for non-essential data collection and surveillance.
 - Informed consent processes that ensure students and families truly understand what they're agreeing to.
 - The right to withdraw consent at any time without penalty.
- **Algorithmic Accountability**
 - Independent audits of algorithms for bias, accuracy, and fairness.
 - Public reporting of outcomes by demographic group to identify disparities.

- Transparent explanations of how algorithmic decisions are made.
- Data Minimization
 - Only collecting data that is truly necessary for educational purposes.
 - Prohibiting the use of student data for marketing or non-educational purposes.
 - Regular deletion of outdated or unnecessary data.
- Privacy by Design
 - Ed-tech tools should be built with strong privacy protections as the default, not as an afterthought.
 - End-to-end encryption and secure data storage.
 - Minimizing the number of people and vendors who can access sensitive information.
- Right to Access and Deletion
 - Students and families should be able to see all data collected about them in a clear, accessible format.
 - The right to request correction or deletion of inaccurate or unwanted data.
 - Mechanisms for appealing or contesting data-driven decisions.
- Human Decision-Making
 - Important decisions (admissions, placement, discipline) should not be made by algorithms alone.
 - Human oversight and robust appeal processes must be in place.
 - Students and families should be able to speak directly with decision-makers.
- Equity Impact Assessments
 - Before adopting new ed-tech, schools should assess potential harms to marginalized students.
 - Ongoing monitoring for disparate impacts on race, class, disability, gender, and immigration status.
 - Involving students and families from diverse backgrounds in decision-making.
- Alternatives to Invasive Proctoring
 - Implementing honor codes, open-note exams, and alternative assessments.
 - Designing assessments that reduce the need for surveillance.
 - Rebuilding cultures of academic integrity based on trust, not suspicion.
- Regulation and Enforcement
 - Strong laws protecting student privacy and prohibiting discriminatory algorithms.
 - Real penalties for violations, including fines and loss of contracts.
 - Independent oversight bodies to investigate complaints and enforce standards.

The Bottom Line

Educational technology is not neutral. Algorithms embed the biases of their creators and the data they're trained on, and surveillance systems treat students as subjects to monitor rather than people to empower. The rush to digitize education has outpaced ethical consideration, regulatory protection, and

evidence of benefit. Students—especially those already marginalized by race, class, disability, or immigration status—bear the greatest burdens and risks. As schools increasingly rely on AI and surveillance tools, we must ask: What kind of education do we want? One that tracks, sorts, and controls? Or one that trusts, nurtures, and liberates? The algorithms we deploy on students today are shaping not just their education, but their relationship to power, privacy, and possibility for the rest of their lives.

References

Chhua, K., Wen, Z., Hathalia, V., Zhu, K. & O'Brien, S. (2024). From Bias to Balance: Detecting Facial Expression Recognition Biases in Large Multimodal Foundation Models. arXiv preprint. <https://doi.org/10.48550/arXiv.2408.14842>

(n.d.). ACLU RESEARCH REPORT.

https://www.aclu.org/wp-content/uploads/publications/digital_dystopia_report_aclu.pdf

Nhando, D. (February 5, 2026). Unmasking EdTech's Surveillance Infrastructure in the Age of AI. TechPolicy.Press. <https://www.techpolicy.press/unmasking-edtechs-surveillance-infrastructure-in-the-age-of-ai/>

(n.d.). FERPA exceptions. Indiana University. <https://ferpa.iu.edu/basics/ferpa-exceptions.html>

(November 30, 2025). FTC Takes Action Against Education Technology Provider for Failing to Secure Students' Personal Data. Federal Trade Commission. <https://www.ftc.gov/news-events/news/press-releases/2025/12/ftc-takes-action-against-education-technology-provider-failing-secure-students-personal-data>

(March 11, 2025). Schools use AI to monitor kids, hoping to prevent violence. Our investigation found security risks. Associated Press. <https://apnews.com/article/25a3946727397951fd42324139aaf70f>

Brislin, S. J., Choi, M., Perkins, E. R., Ahonen, L., McCoy, H., Boxer, P., Clark, D. B., Jackson, D. B. & Hicks, B. M. (2024). Racial Bias in School Discipline and Police Contact: Evidence From the Adolescent Brain Cognitive Development Social Development (ABCD-SD) Study. *Journal of the American Academy of Child & Adolescent Psychiatry* 63(12), pp. 1225-1238. <https://doi.org/10.1016/j.jaac.2024.01.018>

Prothero, A. (February 13, 2024). A New Digital Divide? Low-Income Students See More Ads in the Tech Their Schools Use. Education Week. <https://www.edweek.org/technology/a-new-digital-divide-low-income-students-see-more-ads-in-the-tech-their-schools-use/2024/02>

(n.d.). Algorithmic Accountability Act of 2025. <https://www.congress.gov/bill/119th-congress/senate-bill/2164/text>

1. What is the social issue?

The social issue is the **proliferation of AI-powered educational technologies that embed, amplify, and obscure systemic bias while conducting unprecedented surveillance of students**. AI-powered learning platforms, remote proctoring software, and predictive analytics systems have become deeply embedded across K-12 and higher education, marketed as objective, efficient, and innovative solutions to educational challenges. However, these technologies carry profound risks: they can magnify existing racial, socioeconomic, gender, and disability-based inequalities through biased algorithms trained on historically discriminatory data; they operate as "black boxes" whose decision-making processes are invisible and unaccountable to students, parents, and even educators; they collect vast amounts of personal, behavioral, biometric, and academic data with minimal transparency or meaningful consent; and they make or influence high-stakes decisions—admissions, academic placements, disciplinary actions, intervention targeting, risk flagging—that profoundly shape students' educational trajectories and life opportunities. The increasing reliance on algorithmic systems shifts power away from human judgment, educator expertise, and relational knowledge of students toward opaque computational processes designed by technology companies motivated by profit, not equity. Students are simultaneously subjected to algorithmic sorting that may disadvantage them and digital surveillance that erodes trust, autonomy, privacy, and the conditions necessary for authentic learning—all while institutions present these tools as neutral, progressive, and beneficial.

2. Why does this issue interest you?

This issue is deeply compelling because it sits at the intersection of technology, power, inequality, and the future of education itself—and because the harms are largely invisible by design. Unlike overt discrimination that can be identified, named, and challenged, algorithmic bias operates behind proprietary code, corporate trade secrets, and a veneer of mathematical objectivity that makes it exceptionally difficult to detect, contest, or remedy. When a human admissions officer or teacher makes a biased decision, there is at least the possibility of accountability, appeal, and correction; when an algorithm makes a biased decision, the bias is laundered through the appearance of computational neutrality—"the algorithm decided" becomes an unassailable authority that forecloses challenge. What makes this particularly urgent is that these systems are being adopted at massive scale, often without meaningful evaluation, public debate, informed consent, or input from the communities most affected. Decisions that shape children's and young adults' entire futures—what classes they can take, whether they're flagged as "at risk," whether their exam is invalidated by proctoring software, whether they're admitted to college—are increasingly made by systems that no one fully understands, that were built on data reflecting historical discrimination, and that are sold by companies whose primary obligation is to shareholders, not students. The speed of adoption far outpaces the development of regulation, oversight, or ethical frameworks—creating a situation where an entire generation of students is being sorted, surveilled, and shaped by technologies whose consequences we haven't begun to fully understand.

3. Who is most affected by this issue?

The most affected groups include:

- **Students of color**, particularly Black and Latino/a students, who are disproportionately harmed by predictive algorithms trained on historically biased data—systems that may flag them as "at risk" of dropping out (triggering deficit-based interventions rather than structural support), assign them to lower academic tracks, or subject them to heightened surveillance and discipline based on patterns reflecting systemic racism rather than individual behavior
- **Low-income students** who face compounding disadvantages: they may lack access to reliable internet, private testing spaces, and quality devices required by remote proctoring and AI platforms; they may live in noisy, shared housing environments that proctoring software flags as "suspicious"; and predictive analytics built on socioeconomic data may encode poverty as risk factor, creating self-fulfilling prophecies of failure
- **Students with disabilities** who are disproportionately harmed by remote proctoring software that flags atypical eye movements, body positioning, assistive technology use, or environmental accommodations as "cheating" indicators—penalizing disability-related behaviors as academic dishonesty
- **Students whose first language is not English** who may be flagged by AI writing tools as engaging in plagiarism or AI-generated writing because their syntax patterns differ from training data norms, or whose accents and speech patterns cause errors in voice-recognition systems
- **LGBTQ+ and gender-nonconforming students** whose legal names, gender markers, or presentation may conflict with biometric verification systems, facial recognition software, or institutional records—causing misidentification, outing, or denial of access
- **Black students and students with darker skin tones** who are disproportionately misidentified or failed by facial recognition and proctoring software due to well-documented racial bias in computer vision systems trained predominantly on lighter-skinned faces
- **Students at under-resourced institutions** (community colleges, HBCUs, tribal colleges, rural schools) where ed-tech adoption may be driven by cost-cutting rather than pedagogical improvement, where IT support is minimal, and where students bear the burden of technological failures
- **K-12 students and minors** who cannot meaningfully consent to data collection, have no understanding of how their behavioral and academic data is being harvested and used, and whose digital profiles—created in childhood—may follow them through their entire educational careers and beyond

4. Is this issue mostly about individual choices, social structures, or both?

This issue is almost entirely about **social structures**, with individual choices playing a negligible role for the students most affected. Structurally, the crisis is produced by:

- The political economy of educational technology:** A multi-billion-dollar ed-tech industry aggressively markets AI tools to educational institutions, often with inflated claims of efficacy and neutrality, driven by venture capital investment and profit motives rather than educational mission or equity commitments. Companies like Pearson, Turnitin, Respondus, Proctorio, and others sell products whose algorithms are proprietary trade secrets—unauditable by the institutions purchasing them, let alone by students subjected to them
- Institutional incentives and austerity:** Schools and universities adopt AI tools to reduce labor costs (replacing human judgment with automated systems), manage enrollment and retention metrics (predictive analytics), demonstrate "innovation" to accreditors and boards, and scale operations during budget cuts—prioritizing efficiency and cost-savings over equity, privacy, and pedagogical quality
- The "black box" problem:** Algorithms making consequential decisions about students' lives are proprietary, opaque, and unaccountable. Students cannot understand, inspect, challenge, or appeal algorithmic decisions because the decision-making logic is hidden behind corporate intellectual property protections. There is no due process for algorithmic determination
- Historical data encoding historical discrimination:** AI systems are trained on historical data that reflects centuries of systemic racism, classism, ableism, and other forms of discrimination. When these systems "learn" from this data, they don't discover neutral patterns—they reproduce and automate historical bias at scale, with computational efficiency and the authority of apparent objectivity
- Regulatory vacuum:** There is virtually no federal regulation governing AI use in education. FERPA (Family Educational Rights and Privacy Act) was written decades before AI and is inadequate for governing algorithmic decision-making, biometric data collection, and predictive profiling. Students have few legal protections against algorithmic discrimination
- Power asymmetries:** Students—particularly minors in K-12—have no meaningful ability to refuse these technologies. Participation is mandatory: you must use the learning platform, you must submit to proctoring software, your data will be collected. "Consent" is coerced by the power differential between institutions and students, rendering it meaningless
- Techno-solutionism in education:** A cultural ideology that frames technology as inherently progressive, neutral, and beneficial—obscuring the political choices, economic interests, and power dynamics embedded in technological systems

Individual students cannot opt out, cannot inspect the algorithms affecting them, cannot appeal opaque automated decisions, and often cannot even know that algorithmic systems are influencing their educational experiences. Framing this as a matter of individual choice—students should "be more

careful with their data" or "adapt to the technology"—fundamentally mischaracterizes a structural problem of institutional power, corporate profit, and regulatory failure.

5. Which library resource seems most useful for this topic?

Academic databases such as **ACM Digital Library (Association for Computing Machinery)**, **JSTOR**, and **ERIC (Education Resources Information Center)** would be the most useful library resources for this topic. The ACM Digital Library contains the premier computer science and information technology research, including peer-reviewed work on algorithmic fairness, AI bias, facial recognition accuracy across demographics, and the ethics of automated decision-making—essential for understanding the technical dimensions of how these systems fail. JSTOR provides access to interdisciplinary journals spanning sociology, education, law, and science and technology studies (STS) where critical analyses of ed-tech, surveillance, data privacy, and the social implications of algorithmic systems are published. ERIC indexes education-specific research on the impact of learning analytics, remote proctoring, adaptive learning platforms, and technology integration on student outcomes and equity. Together, these databases span the technical, sociological, legal, and educational dimensions necessary for a comprehensive analysis of this issue.

6. What kind of information might sociologists study about this issue?

Sociologists might study:

- **Science and technology studies (STS) and sociology of algorithms:** How algorithms are socially constructed—built by particular people with particular assumptions, trained on particular data, deployed in particular contexts—and how the myth of algorithmic objectivity obscures the human choices, biases, and power dynamics embedded in these systems
- **Surveillance studies:** How remote proctoring, learning management system tracking, behavioral analytics, and biometric data collection constitute new forms of institutional surveillance that fundamentally alter the student-institution relationship—creating what some scholars call "surveillance pedagogy" where learning environments become sites of constant monitoring
- **Stratification and digital inequality:** How the harms and benefits of educational AI are distributed along lines of race, class, gender, disability, and institutional type—whether algorithmic systems reproduce, amplify, or occasionally disrupt existing educational inequalities, and how the "digital divide" extends beyond access to encompass algorithmic treatment
- **Institutional decision-making and organizational sociology:** How schools and universities decide to adopt AI tools—examining procurement processes, vendor relationships, administrative incentives, the role of marketing claims vs. evidence, faculty and student input

(or lack thereof), and how institutional logics of efficiency and innovation override concerns about equity and privacy

- Race and technology:** How facial recognition systems, proctoring software, and predictive analytics encode and operationalize racial bias—and how the harms fall disproportionately on Black, Indigenous, and Latino/a students through misidentification, differential surveillance, and biased risk predictions

- Power, consent, and agency:** How mandatory technology use in educational settings constitutes coerced participation—students cannot meaningfully consent when refusal means academic penalty or institutional exclusion—and how this reshapes the power dynamics between students, educators, and institutions

- The political economy of ed-tech:** How venture capital, corporate profit motives, and market competition drive the development and adoption of educational AI—examining who profits, who bears the risks, how companies lobby for adoption, and the revolving door between ed-tech companies, policy organizations, and educational institutions

- Data governance and privacy:** How student data is collected, stored, shared, sold, and potentially breached—examining data flows between institutions and third-party vendors, the adequacy of existing privacy frameworks (FERPA, COPPA), and students' understanding of (and control over) their own data

- Resistance and activism:** How students, faculty, parents, and advocacy organizations challenge algorithmic systems—examining protests against proctoring software, campaigns for transparency and accountability, legislative efforts, and alternative models that prioritize human judgment and equity

- Comparative and historical analysis:** How current AI deployment in education compares to previous waves of educational technology adoption (standardized testing, computer-assisted instruction, MOOCs), what historical patterns of techno-enthusiasm and disillusionment reveal, and how different countries regulate educational AI

7. What is one question you would like to ask about this issue?

To what extent do AI-powered predictive analytics systems in education create self-fulfilling prophecies—where students algorithmically flagged as "at risk" based on demographic and socioeconomic data (reflecting historical discrimination rather than individual potential) receive deficit-based interventions, reduced expectations, and constrained opportunities that actually produce the very outcomes the algorithm predicted, thereby automating and legitimizing systemic inequality under the guise of data-driven student support?